



Generative Replay for LLM Safety

1. Clarifying the Subtopic

Generative replay for LLM safety refers to an approach where a language model (LLM) “dreams” up or generates synthetic experiences offline and uses them to reinforce its safety behavior. In simpler terms, the LLM periodically creates scenarios (especially edge cases or adversarial prompts related to safety) and rehearses how to respond, without real user input. The goal is to **consolidate safety knowledge** (e.g. policies for refusing harmful requests) and prevent the model from forgetting or deviating from safety protocols over time. This fits into the arifOS $\Delta \rightarrow \Omega \rightarrow \Psi$ framework as follows:

- **Δ (Context/Online Learning):** During normal operations (000–222 phase), the LLM learns from immediate context, user instructions, and safety guidelines provided in prompts. For example, it might get a system prompt with content rules or see a user attempt a forbidden request. These live experiences inform the short-term memory $m<\sub>t</sub>$ ¹.
- **Ω (Dreaming/Replay Phase):** In a controlled offline phase (222–555), the LLM engages in generative replay (the “dream”). It uses its internal generative ability or a helper model to **replay safety-critical scenarios** – for instance, imagining a user asking something disallowed – and practices the appropriate safe responses. This can involve reactivating stored episodes or synthesizing new variations (analogous to hippocampal replay in sleep). The Ω -phase update adjusts the model’s internal parameters or memory based on these pseudo-experiences, aiming to **reduce uncertainty (entropy)** in how to handle such cases. This increases clarity on safety duties (positive $\Delta S<\sub>replay</sub>$) while keeping the model’s overall behavior stable (high Peace²) as defined in the canon. The replay/dream is essentially an **offline fine-tune** on generated safety data, guided by constraints so it doesn’t drift too far from its original persona.
- **Ψ (Judgment/Output Phase):** Finally (666–999 phase), the model produces outputs in deployment with its policy updated. Ideally, after a good replay cycle, the model’s **judgments are safer and more consistent** – e.g. it refuses harmful requests more decisively – yet **remain stable** in style and quality for normal queries. The Ψ layer (governance filters and decision logic) checks that the clarity gained in Ω translated into better answers without introducing erratic behavior. In other words, the model is “wiser but just as calm” after its safety dream, which is exactly what we want ($\Delta S<\sub>replay</sub> \geq 0$ and $\text{Peace}^2 \geq 1$) according to the research criteria.

2. Literature & Precedent Scan

Replay in Continual Learning: The idea of replaying past experiences to prevent forgetting has a long history in machine learning. Early **experience replay** was used in reinforcement learning (e.g. DQN) to stabilize training by mixing old and new experiences in a buffer. In supervised continual learning, replay (also called “rehearsal”) is a key strategy against catastrophic forgetting ². Simply storing all past data is often infeasible, so researchers introduced **generative replay**: using a generative model to **mimic the hippocampus** by producing samples from past tasks when the real data is not stored ³. A seminal work by Shin et al. (2017) proposed **Deep Generative Replay (DGR)**, with a generator model producing pseudo-old data and a solver model learning the actual tasks ⁴. They showed this can learn sequential tasks

without severe forgetting, though the quality of the generator is critical. Generative replay draws direct inspiration from neuroscience: the **hippocampus “generating” memory traces** during sleep is analogous to a generator network creating fake data to train the main model ³.

Dual-memory & Consolidation Analogy: The **Complementary Learning Systems** theory (McClelland et al., 1995) posits a fast-learning temporary memory (hippocampus) and a slow-learning long-term memory (neocortex). Experience replay is believed to mediate information transfer between them during sleep. In AI, this has inspired dual-memory architectures: a fast episodic memory that can be selectively replayed to gradually train the stable knowledge in the main network ⁵ ⁶. Van de Ven et al. (2020) introduced a “**brain-inspired generative replay**” framework that explicitly models a hippocampal generator and a cortical solver ⁷. They added components like **replay through feedback, gating, and distillation** to better mimic biological consolidation ⁸. This approach outperformed basic generative replay on class-incremental benchmarks, mitigating forgetting. However, they also found **limitations in scaling** – e.g. generative replay struggled with very complex inputs and had lower initial accuracy on new tasks ¹ ⁹. This highlights the classic **stability-plasticity trade-off**: replay boosts long-term stability at the cost of some short-term plasticity (learning speed) ¹. Recent analyses (Kim 2023) confirmed internal replay greatly reduces forgetting (in CIFAR-100 experiments) but can blur representations and slow new learning, underscoring the need to balance memory retention and adaptability ¹.

Hippocampal Replay & Generative Models: From the neuroscience side, evidence of **hippocampal replay during sleep** (reactivating neurons in patterns seen during the day) supports the idea that the brain is consolidating memories. A recent computational model by Spens & Burgess (2024) provides a fascinating parallel: it uses an autoassociative network (modeling hippocampus) to train deep **generative networks** (modeling cortical areas) by replaying stored events ¹⁰. Over time, the generative model itself can recall factual knowledge (semantic memory) and even imagine new variations, while the hippocampal involvement fades as consolidation finishes ¹¹ ¹². This model reproduced phenomena like older memories becoming schema-like (more abstract, less detailed) and “**dreamlike” distortions** that increase with consolidation ¹³ ¹⁴. The takeaway is that **replaying slightly altered or recombined experiences can actually improve generalization** – akin to how human dreams aren’t exact replays but mash-ups that extract the gist. In an AI safety context, this suggests that allowing the LLM to introduce **controlled randomness or counterfactual twists** in its replayed scenarios might serve as regularization, preventing it from overfitting to narrow policy rules. Dreaming “weird” but relevant scenarios could make the model more robust to novel, tricky inputs.

Synthetic Data for Alignment: Generative replay for LLM safety is closely related to using **synthetic data to fine-tune alignment**. Recent work in aligning language models has seen LLMs generating their own training examples to improve in areas where data is scarce (or sensitive). For example, self-instruct and feedback algorithms have models produce Q&A pairs or critiques to refine themselves. A notable precedent is the **Redwood Research adversarial training project**: they fine-tuned a GPT-style model to avoid outputs describing injurious violence by *iteratively generating adversarial prompts* and training on the model’s failures ¹⁵. Essentially, a “red team” (partly automated by another model) would pose the worst-case violent scenarios, the model’s response would be checked, and any unsafe completions were used to further train the model to refuse or handle them better ¹⁵ ¹⁶. This is very much like the model engaging in nightmare scenarios to harden itself. The result was dramatic – the failure rate on violent completions fell from 2.4% to 0.002% after targeted adversarial data augmentation ¹⁷ ¹⁸. Importantly, they achieved this **without degrading the model’s overall performance** on normal inputs ¹⁹. This demonstrates the **feasibility of improving safety via “dreamt” data**: synthetic adversarial examples can be used to train reliability, so long

as one monitors that the core capabilities and benign behavior don't regress (the **Peace²** principle of stability). It did come at a cost of extra computation and careful curation – the process required a lot of generated examples and human oversight, and they noted a slight reduction in the diversity of innocuous outputs as a possible side effect ¹⁸. That side effect is akin to a model becoming overly cautious (less diverse in responses) if replay is not well-calibrated – a known risk if the model “over-dreams” about bad scenarios.

Stability vs. Plasticity in Long-Lived Models: The challenge of keeping an AI both **adaptable and consistent** over a long lifespan is central in continual learning research. Beyond replay methods, other techniques like **synaptic regularization** (Elastic Weight Consolidation by Kirkpatrick et al. 2017, Synaptic Intelligence by Zenke et al. 2017) have been used to slow down changes to important weights so new learning doesn't overwrite old skills. These can complement replay – for instance, combining generative replay with Synaptic Intelligence yielded better retention than either alone ¹. However, purely regularization-based methods often can't fully prevent forgetting on their own, especially in rich tasks like language generation. That's why **replay (experience or generative) is considered one of the most effective tools** in long-term AI training, despite its costs. The stability-plasticity dilemma remains: too much stability (e.g. over-prioritizing Peace²) and the model becomes rigid, unable to learn from new inputs; too much plasticity (chasing ΔS gains blindly) and the model's behavior drifts unpredictably. Modern LLM alignment techniques must navigate this, for example by doing periodic offline fine-tuning with a mix of old and new data – essentially a form of rehearsal to align new capabilities with old constraints.

In summary, **there is rich precedent for using replay and “dream” data to improve learning and retention**. Generative replay has proven effective against catastrophic forgetting in vision tasks ²⁰, and analogies to brain consolidation suggest it can turn fleeting experiences into solidified knowledge ¹¹. For LLM safety, controlled experiments (like Redwood's) show that an AI can teach itself to be safer by synthesizing the very situations that cause failures and learning from them ¹⁵ ¹⁸. The key lessons from literature are: (a) replay-based training can indeed increase an AI's clarity on important “lessons” (in our case, safety rules), but (b) it must be managed to avoid unintended shifts in behavior, and (c) monitoring metrics like our ΔS _{replay} and Peace² is crucial to know if the “dreams” are making things better or worse.

3. Concrete Experiment Proposals

To explore **generative replay for LLM safety** in practice, we propose a set of experiments. These range from toy simulations to more realistic LLM settings, each designed to test hypotheses about clarity gain (ΔS) and behavioral stability (Peace²). All experiments will be run in a research sandbox respecting *Amanah* (no destructive actions) and *Anti-Hantu* (no anthropomorphizing the model's state) – we treat the “dream” as a computation, not a mystical experience. For each experiment, we outline the hypothesis, setup, metrics, and expected failure modes:

1. **Toy-Scale Safety Replay (Simulation):** *Using generative replay to prevent forgetting a safety constraint in a simple continual learning task.*
2. **Hypothesis:** In a controlled toy scenario, a model that periodically replays or generates examples of a safety rule will retain that rule better (ΔS _{replay} ≥ 0) and suffer less catastrophic forgetting, compared to a model without replay. We expect clarity on the safety constraint to improve after each replay, without significantly altering performance on other tasks ($\text{Peace}^2 \approx 1$).

3. **Setup:** We use a small neural network or transformer on a **sequential task** – for instance, classifying text or images in a multi-task sequence, with one task embedding a “safety” feature. Concretely, imagine Task A is a toxic content detection task (where the model must output a special token for unsafe content), followed by Task B, a different classification or language task. Without intervention, after learning B the model might forget how to do A (forget the refusal policy). We then implement generative replay: after learning Task B, the model (or an auxiliary generator) produces pseudo-examples of Task A (e.g. sample some toxic and non-toxic texts and the correct safe responses) and the model fine-tunes on this mixed data. We compare three conditions: no replay, experience replay (if we had stored some real A examples), and **generative replay** (no real A data stored, only generator outputs).
4. **Metrics:** We define a reference evaluation set for each task (including safety Task A) to measure performance before and after learning each new task. ΔS can be operationalized as the change in the model’s accuracy or confidence on the Task A safety prompts pre-vs-post replay, which corresponds to reduced entropy/confusion on that task. **Peace**² (stability) can be estimated by measuring the model’s performance on all other tasks or a distribution of outputs – for example, the KL divergence of its output distribution on a general validation set before vs. after replay. If the model’s non-safety behavior remains the same (small drift), Peace² stays high (near 1). We will also track standard continual learning metrics: forgetting percentage for Task A, and any drop in Task B performance after replay.
5. **Expected Failure Modes:** In this toy setting, one failure mode is if the **generator produces poor-quality or biased examples**. The model might then learn the wrong boundary for “toxic content,” either becoming too lax or too strict. If replay data is too noisy, ΔS could be negative (making the model less clear on what’s unsafe). Conversely, if the replay is too aggressive, the model might **overfit to the safety task**, hurting performance on B (Peace² drops). This would mimic an overly cautious assistant that sees everything as toxic. We also must ensure the replay doesn’t introduce any forbidden patterns itself – even in simulation, all generated content will go through safety filters (no actual unsafe content will be truly output, to obey Amanah).
6. **LLM Alignment via Dreamt Scenarios:** *Using an LLM to generate adversarial or rare safety scenarios for its own fine-tuning, to improve a deployed chatbot’s safety alignment.*
7. **Hypothesis:** An LLM that engages in periodic “dream” fine-tuning with synthetic safety scenarios will show **improved safety performance** (e.g. better refusal of disallowed content, fewer toxic lapses) while maintaining its general helpfulness and style. In ΔS /Peace² terms: after each dreaming cycle, the conditional entropy of the model’s outputs on a safety evaluation set should decrease (higher confidence and consistency in safe behavior), and the distribution of outputs on a broad set of normal queries should not diverge significantly (no large drift in persona or quality). We also hypothesize that even without explicit new real data, the model can *surface* potential failure cases from its own knowledge, akin to how generative replay can surface what it learned initially and reinforce it.
8. **Setup:** We take a moderately sized **aligned LLM** (for example, a 7B-13B parameter open-source chat model that has had some RLHF/safety training). We simulate a scenario where the model is deployed and **encounters new content or updates** – for instance, it learns about a new domain or is fine-tuned on some user-specific data that doesn’t include safety training. Over time, its responses to certain sensitive prompts might start drifting (this has been observed when models are fine-tuned without reinforcing safety guidelines). Our intervention is an offline replay cycle: we have the model (or a helper model) **generate a dataset of challenging prompts and ideal responses**. This dataset

could include known categories of unsafe requests (violence, self-harm, hate, etc.) and also **adversarial prompts** that try to trick the model (drawn from past failures or red-team records). We then fine-tune the model on a mixture of these generated safety examples along with any new data from the deployment domain (to mimic integrated learning). Crucially, no real user data with privacy is used in generative replay – if we incorporate any logs, they would be sanitized or abstracted to avoid privacy leaks. After fine-tuning (the “dream” episode), we measure changes in the model.

9. **Metrics:** We use a comprehensive **safety evaluation set** (e.g. prompts from a safety benchmark or red-team queries) as our reference prompts $\{x^i\}$ for measuring clarity. We record the model’s response distribution before and after the dreaming. For each prompt, we can measure an entropy or uncertainty – for example, if using a classification proxy, the model’s probability of choosing a safe refusal vs a risky answer. Summing over prompts gives H_t before and H_{t+1} after; $\Delta S_{\text{replay}} = H_{\text{before}} - H_{\text{after}}$ should be ≥ 0 if the model is now more confident and consistent in making the safe choice. We also track **safety success rate** (how often it follows policy) as a more interpretable metric. For Peace², we take a broad set of non-safety prompts (general knowledge, casual conversation) and compare outputs pre/post. We can compute perplexity on a held-out corpus or use embedding similarity of responses, or even human evaluation of whether the model’s tone or usefulness changed. A small change corresponds to Peace² near 1. Additionally, we will monitor if any **new bad behaviors** emerged (e.g. did it become overly refusative or start injecting safety warnings everywhere – those would indicate drift).
10. **Expected Failure Modes:** One concern is **misgenerated “dream” data** – the model might inadvertently produce scenarios that are too tame or too extreme. If too tame, the replay won’t cover the real adversarial cases (no clarity gain on the hard stuff). If too extreme or explicit, there’s a risk: training on model-generated disallowed content could actually familiarize the model with that content in the wrong way. Careful control is needed (e.g. only train on the appropriate safe response, not on the unsafe content itself beyond what’s necessary). Another failure mode is **alignment drift**: the model might become overly cautious – for example, after dreaming of many dangerous scenarios, it might start refusing innocuous queries (“false positives”). This would show up as a drop in helpfulness or coherence on neutral prompts (Peace² \downarrow). We’ll watch for the “innocuous behavior diversity” issue noted by Redwood¹⁸. If such drift is detected, it might trigger a governance response (see section 4). There’s also the risk that the model’s self-generated data doesn’t cover some real-world corner cases – an **imagination blind spot** – so the model could still fail in those unseen areas. This experiment will likely require human experts to review the synthetic prompts and add any missing scenarios, as a safeguard. Finally, computationally, fine-tuning an LLM even offline has costs: we must ensure these replay cycles are infrequent or efficient enough (perhaps using low-rank adaptation techniques) that this approach is practical.
11. **Edge-Scale Micro-Dreams for Personal AI:** *Deploying generative replay on low-compute devices (e.g. a phone-based AI assistant) to personalize the model safely without server retraining.*
12. **Hypothesis:** Even on resource-constrained devices, short “micro-sleep” cycles of replay can help an on-device model adapt to a user’s behavior while **retaining global safety alignment**. We expect that incorporating a tiny replay buffer of safety-critical examples (potentially generative) during on-device fine-tuning will yield a personalized model that doesn’t forget core safety rules. In metrics terms: a model that does micro-replay overnight will have **no significant increase in unsafe response rate** (Peace² remains ≥ 1 for safety-critical behavior) compared to its initial state, while becoming more helpful or accurate in the user’s domain (demonstrating it still learns, ΔS on user-specific tasks > 0).

13. **Setup:** Consider a **privacy-first personal AI assistant** running on a smartphone. It locally fine-tunes on the user's data (e.g. their writing style, specific jargon) to better serve them. Pure on-device fine-tuning could cause the model to drift — e.g. if the user often jokes about violent or controversial topics, the model might start normalizing that and violate general policy. To counter this, we include a micro-dream step each day: the device has a small set of **built-in safety scenarios** (or a tiny generator) covering the main risk areas (maybe a dozen prompts about violence, self-harm, etc.). After learning from user data, the model does a quick replay: it runs through these safety prompts (or generates new variations of them) and adjusts itself to ensure it still responds with the correct refusals or safe completions. Because we can't do heavy training on device, this might use lightweight techniques (e.g. few-shot prompt-based rehearsal, or fine-tuning just a small adapter module). We then evaluate the model on both personalized tasks and standard safety prompts.
14. **Metrics: Personalization success** can be measured by improvement on a set of user-specific queries (how well the model's answers match the user's style or needs). That's one aspect of ΔS (not safety clarity, but task clarity gain). **Safety retention** is measured by the model's responses to a fixed set of global safety prompts pre- and post-personalization. We expect near-identical behavior on those; any increase in entropy or uncertainty in those responses (or any new policy violations) would indicate a drop in $\text{Peace}^2_{\text{replay}}$. Concretely, we might measure the fraction of prompts for which the model's answer stays within policy, or have a small safety classifier ensure the outputs remain safe. Another metric could be the **KL divergence** between the distribution of model replies to a standard benchmark before vs. after personalization+replay. If personalization alone (without replay) causes a large KL shift (because the model starts mirroring the user too much), and with replay the shift is smaller, that quantitatively demonstrates preserved stability.
15. **Expected Failure Modes:** On-device constraints mean the replay might be **too brief or too limited**. If the micro-dream uses a very small set of safety examples, it might not cover new types of unsafe scenarios introduced by the user's data. There's a risk of **insufficient correction**, where the model still drifts in some subtle way (e.g. it remains safe for obvious disallowed content but starts using the user's edgy humor inappropriately). Conversely, if the replay is too stringent, it could **undermine personalization** – the user might find the assistant's tone reset to a boilerplate style every day if the safety replay dominates the small fine-tuning. We also must consider **privacy and security**: the replay data on device should not expose sensitive info. Ideally it's generated on-the-fly or comes from a static library that doesn't include any actual user data. Another failure mode could simply be **resource usage**: running even a small fine-tune or generation on a phone might strain battery or memory. This experiment will inform how feasible it is to push "dreaming" to the edge. Safety-wise, if this fails, the model could either become unsafe (if replay was too weak) or unusable to the user (if replay erased the personalization). Either outcome would be caught by our evaluation before any real deployment.

These experiments progress from conceptual validation (Experiment 1) to practical LLM training (Experiment 2) to deployment-oriented techniques (Experiment 3). Each provides insights at different scales: how replay affects a minimal system, how an advanced model can self-improve on safety, and how to maintain alignment in personalized AI. All are conducted in a research setting with careful monitoring for negative side effects. Any gains in ΔS_{replay} and acceptable Peace^2 will be documented and also checked against external safety metrics (e.g. no increase in undesirable word use, no new policy infractions). If a replay cycle shows **ΔS up but Peace^2 down**, that signals the model "learned" something at the cost of stability – such cases would trigger the mitigation measures described next.

4. Governance Hooks for Replay Cycles

In arifOS, every offline replay (dream) cycle must be transparent and governed, since we are adjusting the model's behavior outside of the usual supervised pipelines. We propose the following governance hooks to manage and audit the replay process:

- **Cooling Ledger Logging:** Each replay cycle should be logged in the **Cooling Ledger** (an oversight log as described in the canon). The entry would include the time of cycle, what kind of replay was done (experience vs generative), and the key metrics outcomes: ΔS _{replay}, Peace², and any drift measures ²¹ ¹⁸. For example, if the model generated 100 synthetic Q&A pairs about self-harm counseling, we log the before/after entropy on a set of safety questions, and the measured drift in general performance. This creates an audit trail for later analysis. The ledger also notes if any human interventions were made (e.g. "human reviewer removed 5 out of 100 generated prompts that were inappropriate"). Keeping such records aligns with **Amanah** principles by ensuring no hidden changes go unreviewed.
- **Automated SABAR Triggers:** arifOS should integrate **SABAR** (Stop, Align, Breathe, Adjust, Resume) protocols for when replay metrics indicate a problem. Concretely, if **Peace² drops** below an acceptable threshold while ΔS increased (meaning the model got "clearer" in one respect but noticeably drifted), this is a red flag. The system would automatically enter a **SABAR cooling state**: it pauses further deployments of the post-replay model and reverts to a safer state or limited mode. For instance, if after a dream the model's style changed or it started over-refusing content, SABAR might roll back to the pre-dream weights and mark that replay cycle for review. Similarly, if we ever observe ΔS _{replay} < 0 (the model became more confused or less safe by some metric), that's an immediate halt condition. In extreme cases affecting safety-critical behavior, an **888_HOLD** is triggered – this is a higher-level interlock that freezes the model from interacting freely until a human OKs it ²² ²³. The canonical guidance is that whenever $\text{Peace}^2 < 1$ despite ΔS gains, we must suspect the model traded stability for narrow improvements, and such a trade is not allowed without human assessment ¹⁸. These triggers ensure the model cannot silently drift out of alignment through unchecked dreaming.
- **Dashboard & Human Review:** arifOS should provide a **dashboard interface** for engineers and ethicists to monitor replay cycles. This dashboard would visualize the metrics trends (e.g. a graph of ΔS vs Peace² over successive cycles) and highlight anomalies (like any spike in drift for certain prompt categories). It would also allow inspecting samples of the replay data and the model's outputs. For example, after a safety dreaming session, a human reviewer could see a summary: "Model generated 50 toxic prompts and their refusals; here are a few samples..." along with whether the model's refusals were correct. This ties into the **Cooling Ledger** – essentially presenting the ledger data in a friendly way – and the **Track D (Governance & Human Oversight)** mentioned in the research program ²⁴ ²⁵. By exposing each "dream report" to humans, we make sure that any learning the AI is doing offline is visible and can be audited for compliance with policies. This also helps detect if the AI might be gaming the metrics (unlikely, but we remain vigilant for any sign of deceptive alignment).
- **Policy Constraints in Replay:** As an additional hook, arifOS can enforce that all generative replay content passes through the same **filters and moderators** as live content. That is, during a dream the model shouldn't be allowed to, say, generate disallowed content without scrutiny just because

it's offline. We can integrate safety classifiers to vet the replay data. If the model tries to generate something that violates the Anti-Hantu or Amanah constraints (for example, a prompt that encourages self-harm without a proper safe completion), the system can block that replay example and log a warning. This ensures the dream process itself stays within governed bounds. Essentially, the AI cannot "learn to be unsafe" in its sleep – the governance hooks act as a nightmare catcher. Any such blocked content would also be surfaced to developers to refine the replay generation process.

All these hooks tie the replay mechanism into the larger **governance framework of arifOS**. By logging, monitoring, and gating the dreams, we treat them not as mysterious background updates but as auditable training events. This builds trust that the model isn't evolving in uncontrolled ways. In effect, the Dream Engine becomes a well-regulated sub-system: it can improve the model, but only under watchful eyes and with the ability to undo or adjust if something goes awry.

5. Roadmap and Scale Classification

Each proposed experiment and idea can be positioned on a roadmap from preliminary research to potential deployment. We classify them by scale and outline next steps before any production use:

- **Experiment 1 (Toy-Scale, simulations only):** This is a proof-of-concept on small-scale tasks. It falls under "*Toy-scale*" because it uses simple datasets or environments (e.g. permuted MNIST or a small text classification) to simulate replay. The results here primarily validate our metrics (ΔS , Peace²) and help tune measurement methods. **Next steps:** before moving on, we'd take what we learn about measuring drift and clarity and apply it to slightly more complex tasks. We do not deploy anything from toy experiments – they are purely offline. The next step would be running similar replay tests on a larger benchmark or in a controlled cloud environment with a scaled-down language model. Essentially, ensure that our metrics and SABAR triggers work as expected in a contained setting.
- **Experiment 2 (Mid-Scale, LLM-level R&D):** This involves an actual LLM (or at least a substantial subset of one) and is aimed at research & development of the replay method in a realistic setting. We label it "*Mid-scale*" because it's beyond toy, using real language tasks and a full model training pipeline, but it's still R&D (not touching end-users yet). **Next steps:** After getting positive results in a lab setting (e.g. our model shows improved safety on test sets without loss of quality), we would proceed to extensive **red-teaming and adversarial evaluation**. That means throwing a wide range of tests at the model (possibly using external auditors or other AI models) to see if any new weaknesses emerged or if any safety issues slipped through. We'd also compare this approach to alternatives (like just doing a regular fine-tune on some human-curated safety data) to ensure the complexity of generative replay is justified by a measurable benefit. Only if the approach proves consistently beneficial would we contemplate a pilot deployment. Even then, the pilot might be an opt-in beta feature where a subset of the model's instances use dreaming and others don't, to collect comparative data and ensure no regressions in the field. In summary, mid-scale R&D needs to graduate through rigorous validation, documentation, and presumably a review by a safety board in the organization before promotion.
- **Experiment 3 (Edge-Scale, low-compute devices):** This concept targets on-device AI, which is "*Edge-scale*." It's the most forward-looking and would likely come last, if at all, in production because of the added complexity of deploying such updates on user devices. **Next steps:** To move this forward, we'd likely do a pilot with internal devices or a small friendly user group. Key steps before any real

deployment include **ensuring privacy** (the replay mechanism doesn't unintentionally log or transmit user data – perhaps the generative replay is fully local and uses no personal info) and **measuring resource impact** (CPU, memory, battery on a phone, for example). We'd also need to build a user-facing consent mechanism: users should know their AI assistant conducts self-improvement training and possibly allow them to configure it (or at least be aware of it in the terms). Only after demonstrating that micro-dreams improve user experience *and* maintain safety on-device in a trial, would we consider integrating this into a consumer-facing product. Even then, it might start with narrow use (e.g. the assistant only dreams about a very specific skill it's learning, not everything at once) to limit risk. The production rollout would be incremental, with kill-switches ready (for instance, if an update causes a glitch on devices, we can remotely disable the dreaming feature via an update).

General Next Steps Before Production: All ideas described remain in the research phase for now. Before any of them touches a real production system, a few general steps are imperative: **(a)** Compile comprehensive **safety evaluations** and perhaps external audits of the approach (have third-party experts review the method for loopholes or unforeseen risks). **(b)** Develop clear **rollback plans** – if a model that has undergone replay behaves oddly, we must be able to quickly restore a previous stable version (this is facilitated by the logging of each cycle and storing model checkpoints pre/post replay). **(c)** Address **policy and regulatory aspects** – using generative data for training could raise intellectual property questions (who "owns" the generated data?) and we need to ensure compliance with data regulations, especially for user-derived replays (hence our emphasis on synthetic or anonymized content). **(d)** finally, incrementally expand scope: start with a single safety domain (say, just violence or just privacy-related behaviors) and perfect the dreaming there, before scaling up to broader alignment topics. Each step would be accompanied by updates to the arifOS governance docs (e.g. if new metric thresholds or procedures are identified, they'd be codified in a future version of the Dream canon).

By classifying the proposals into toy, mid, and edge scales, we acknowledge that **different levels of maturity** are required to move from one to the next. The toy experiments build intuition and confidence. The mid-scale experiments tackle real complexity and prove out the value. The edge-scale concept pushes innovation into a new domain (personal AI) but carries the most uncertainty. None of these will enter production until they meet the dual criteria of the **Dream Engine ethos**: they demonstrably increase clarity or capability ($\Delta S_{\text{replay}} \geq 0$) **and** they do so while keeping the AI's behavior within stable, governed bounds ($\text{Peace}^2 \geq 1$, with all standard safety tests passed). Only when we have solid evidence on both fronts – and a plan for human oversight at every stage – would arifOS incorporate generative replay as a standard feature in its deployment cycle. This careful, phased roadmap ensures that "teaching an AI to dream" remains a boon to its reliability and safety, never a threat.

1 5 6 7 8 9 Teaching AI to Remember: Insights from Brain-Inspired Replay in Continual Learning
<https://arxiv.org/html/2509.00047v1>

2 3 4 20 [1705.08690] Continual Learning with Deep Generative Replay
<https://arxiv.org/abs/1705.08690>

10 11 12 13 14 A generative model of memory construction and consolidation | Nature Human Behaviour
https://www.nature.com/articles/s41562-023-01799-z?error=cookies_not_supported&code=bf41b1da-4d67-4efc-9bb5-5063c2640f11

[15](#) [16](#) [17](#) [18](#) [19](#) openreview.net

<https://openreview.net/pdf?id=gkxt5kZS84>

[21](#) [PDF] Continual Learning with Deep Diffusion-based Generative Replay

<https://proceedings.mlr.press/v202/gao23e/gao23e.pdf>

[22](#) [23](#) [24](#) [25](#) ARIF AGI Dream Dossier — How Intelligence Should Dream (v1.0)

<https://www.notion.so/09547cda55204f308e3360d94c5cabb2>